

ai_hardware_hub_v1.json

```
[
  {
    "id": "GPU-NV-H100",
    "name": "NVIDIA H100",
    "manufacturer": "NVIDIA",
    "architecture": "Hopper",
    "memory_size_gb": 80,
    "memory_type": "HBM3",
    "memory_bandwidth_tbs": 3.35,
    "fp8_tflops": 3958,
    "fp16_tflops": 1979,
    "tdp_watts": 700,
    "release_year": 2022,
    "use_case": "Training & Inference",
    "interconnect": "NVLink 4.0"
  },
  {
    "id": "GPU-NV-B200",
    "name": "NVIDIA B200 (Blackwell)",
    "manufacturer": "NVIDIA",
    "architecture": "Blackwell",
    "memory_size_gb": 192,
    "memory_type": "HBM3e",
    "memory_bandwidth_tbs": 8.0,
    "fp8_tflops": 9000,
    "fp16_tflops": 4500,
    "tdp_watts": 1000,
    "release_year": 2024,
    "use_case": "Next-gen LLM Training & Inference",
    "interconnect": "NVLink 5.0"
  },
  {
    "id": "GPU-AMD-MI300X",
    "name": "AMD Instinct MI300X",
    "manufacturer": "AMD",
    "architecture": "CDNA 3",
    "memory_size_gb": 192,
    "memory_type": "HBM3",
    "memory_bandwidth_tbs": 5.3,
    "fp8_tflops": 2600,
    "fp16_tflops": 1300,
    "tdp_watts": 750,
    "release_year": 2023,
    "use_case": "Large Model Inference & Memory-intensive Workloads",
    "interconnect": "Infinity Fabric"
  },
  {
    "id": "ACC-INT-G3",
    "name": "Intel Gaudi 3",
    "manufacturer": "Intel",
    "architecture": "Gaudi",
    "memory_size_gb": 128
  }
]
```

```
"memory_type": "HBM2e",  
"memory_bandwidth_tbs": 3.7,  
"fp8_tflops": 1800,  
"fp16_tflops": 1800,  
"tdp_watts": 900,  
"release_year": 2024,  
"use_case": "Cost-effective AI Training & Inference",  
"interconnect": "Ethernet (RDMA)"
```

```
}
```

```
]
```