

ai_hardware_hub_v2json

```
[
  {
    "id": "GPU-NV-H100",
    "name": "NVIDIA H100",
    "manufacturer": "NVIDIA",
    "architecture": "Hopper",
    "memory_size_gb": 80,
    "memory_type": "HBM3",
    "memory_bandwidth_tbs": 3.35,
    "fp8_tflops": 3958,
    "fp16_tflops": 1979,
    "tdp_watts": 700,
    "release_year": 2022,
    "use_case": "Training & Inference",
    "interconnect": "NVLink 4.0"
  },
  {
    "id": "GPU-NV-B200",
    "name": "NVIDIA B200 (Blackwell)",
    "manufacturer": "NVIDIA",
    "architecture": "Blackwell",
    "memory_size_gb": 192,
    "memory_type": "HBM3e",
    "memory_bandwidth_tbs": 8.0,
    "fp8_tflops": 9000,
    "fp16_tflops": 4500,
    "tdp_watts": 1000,
    "release_year": 2024,
    "use_case": "Next-gen LLM Training & Inference",
    "interconnect": "NVLink 5.0"
  },
  {
    "id": "GPU-AMD-MI300X",
    "name": "AMD Instinct MI300X",
    "manufacturer": "AMD",
    "architecture": "CDNA 3",
    "memory_size_gb": 192,
    "memory_type": "HBM3",
    "memory_bandwidth_tbs": 5.3,
    "fp8_tflops": 2600,
    "fp16_tflops": 1300,
    "tdp_watts": 750,
    "release_year": 2023,
    "use_case": "Large Model Inference & Memory-intensive Workloads",
    "interconnect": "Infinity Fabric"
  },
  {
    "id": "TPU-GO-V5P",
    "name": "Google TPU v5p",
    "manufacturer": "Google",
    "architecture": "TPU v5",
    "memory_size_gb": 95,
```

```

        "memory_type": "HBM",
        "memory_bandwidth_tbs": 4.8,
        "bf16_tflops": 459,
        "int8_tops": 918,
        "tdp_watts": 450,
        "release_year": 2023,
        "use_case": "Large-scale Training (Google Cloud)",
        "interconnect": "ICI v3"
    },
    {
        "id": "LPU-GR-LPU1",
        "name": "Groq LPU (Language Processing Unit)",
        "manufacturer": "Groq",
        "architecture": "TSP (Tensor Streaming Processor)",
        "memory_size_gb": 0.23,
        "memory_type": "SRAM",
        "memory_bandwidth_tbs": 80.0,
        "int8_tops": 750,
        "fp16_tflops": 188,
        "tdp_watts": 300,
        "release_year": 2023,
        "use_case": "Ultra-low Latency Inference",
        "interconnect": "GroqLink"
    },
    {
        "id": "ACC-AWS-TR2",
        "name": "AWS Trainium2",
        "manufacturer": "AWS",
        "architecture": "Trainium",
        "memory_size_gb": 96,
        "memory_type": "HBM3",
        "memory_bandwidth_tbs": 4.0,
        "fp8_tflops": 1300,
        "tdp_watts": 600,
        "release_year": 2024,
        "use_case": "High-efficiency Training (AWS)",
        "interconnect": "EFA"
    },
    {
        "id": "NPU-AP-M4U",
        "name": "Apple M4 Ultra (Neural Engine)",
        "manufacturer": "Apple",
        "architecture": "M4",
        "npu_ops_tops": 38,
        "memory_size_gb_shared": 192,
        "memory_bandwidth_tbs": 0.8,
        "tdp_watts": 120,
        "release_year": 2025,
        "use_case": "On-device / Edge AI Development",
        "interconnect": "UltraFusion"
    }
]

```