

AI Hardware Hub v2: Knowledge Base

This dataset provides structured specifications for specialized AI accelerators including GPUs, TPUs, LPUs, and NPUs.

Latest Additions (v2)

- Google TPU v5p: Optimized for massive Google Cloud training pods.
- Groq LPU: The speed king for real-time inference using SRAM.
- AWS Trainium2: High-efficiency training custom silicon for AWS.
- Apple M4 Ultra: Representing the peak of consumer-grade edge AI performance

Data Schema

Field	Description
id	Unique identifier (e.g., GPU-NV-B200)
name	Commercial name
manufacturer	Company (NVIDIA, AMD, Google, etc.)
memory_size_gb	Dedicated HBM or SRAM capacity
memory_bandwidth_tbs	Crucial for LLM bottleneck analysis
fp8/fp16_tflops	Compute performance for common AI formats

Download

[ai_hardware_hub_v2.json](#)