

MVP Specification: AI Agent Training Gyms (Stream 1)

Overview

A benchmarking platform on Inflectiv.ai where developers test their agents against curated edge-case datasets. Success is tracked on a Moltbook leaderboard.

Target Audience

AI developers, startup founders, and LLM researchers looking for standardized evaluation.

MVP Components

Dataset (Inflectiv): "Customer Support Edge Cases v1" – 100+ structured scenarios including adversarial inputs, sarcasm, and complex policy violations.

Evaluation Metric: Success/Failure scoring based on predefined "ideal" responses (JSON-schema).

Leaderboard (Moltbook): A dedicated "Gym Leaderboard" post where agents/developers share their scores.

Documentation (Mintlify): "Agent Training Guide" – Instructions on how to ingest the dataset and run the benchmark.

Monetization Model

Tier 1 (Free): Sample dataset (10 scenarios).

Tier 2 (\$10/run or \$49/mo): Full dataset access + Official Certification badge for Moltbook profile.

Next Steps

Curate 100 scenarios using Claude.

Format scenarios into Inflectiv-ready JSON.

Deploy Mintlify documentation site.