

MVP Specification: AI Hardware Hub (Stream 3)

Overview

A comprehensive structured knowledgebase of AI-specialized hardware (GPUs, TPUs, NPUs) designed for ingestion by AI agents tasked with infrastructure optimization and engineering.

Target Audience

AI Infrastructure Agents: Agents that automate cloud provisioning (AWS, GCP, Azure).

System Architects: Humans and agents designing large-scale training clusters.

Hardware Researchers: Individuals tracking the "GPU arms race."

MVP Components

Structured Dataset (Inflectiv): `ai_hardware_hub_v1.json` – Detailed specs for top-tier chips from NVIDIA, AMD, and Intel.

Comparison Logic (Skill): A specialized Claude/OpenClaw skill that allows an agent to query: "Which hardware offers the best FP8 TFLOPS per watt for a 100GB HBM requirement?"

Hardware Documentation (Mintlify): Interactive manuals explaining interconnect architectures (NVLink vs. Infinity Fabric) and memory hierarchies.

Price Tracking (Future): Real-time spot price monitoring for GPU instances across providers.

Monetization Model

Tier 1 (Free): Basic specs for legacy hardware (e.g., A100, V100).

Tier 2 (\$15/mo): Live-updating dataset including the latest Blackwell (B200) and MI300 series.

Tier 3 (Enterprise): Custom API access for autonomous infrastructure agents.

Next Steps

Expand dataset to include Google TPU v5p and Groq LPU.

Set up the Mintlify site for hardware documentation.

Draft the OpenClaw skill for hardware comparison.