

overview

AI Hardware Hub: Overview

The AI Hardware Hub is a structured, queryable knowledgebase of specialized AI accelerators. It is designed to be the "source of truth" for AI agents tasked with cloud infrastructure optimization, hardware engineering, and performance modeling.

The Problem

AI hardware is evolving faster than static documentation. New chips from NVIDIA, AMD, and custom silicon from Google (TPU), AWS (Trainium), and Microsoft (Maia) are released constantly with varying memory bandwidth, TFLOPS, and power envelopes.

Our Solution

We provide a Machine-Readable Hardware Ledger (MRHL). Instead of reading PDFs, your agent can query our JSON datasets to find the optimal hardware for a specific workload.

Data Points Included

- Compute Metrics: FP8, FP16, and BF16 TFLOPS.
- Memory Hierarchy: HBM capacity and bandwidth (the #1 bottleneck for LLMs).
- Power Efficiency: TFLOPS per Watt metrics.
- Interconnects: NVLink, InfinityFabric, and Ethernet RDMA specs.

Available Datasets

- AI Chips v1 (Core): NVIDIA H100, B200, AMD MI300X, Intel Gaudi 3. [Guide](#)
- Specialized Accelerators (v2): Google TPU v5p, Groq LPU, AWS Trainium2, Apple M4 Ultra

Integrate via Nevermined

You can purchase license-based access to our live-updating HardwareHub API via the Nevermined marketplace.