

SKILL

name: ai-hardware-advisor description: Expert system for selecting and optimizing AI specialized hardware (GPUs, TPUs, LPUs). version: 1.0.0 author: SymbioStreams tags: [hardware, ai-infrastructure, optimization, silicon]

AI Hardware Advisor

The AI Hardware Advisor is a specialized skill for Claude and other AI agents. It leverages the SymbioStreams AI Hardware Hub to provide data-driven recommendations for AI infrastructure.

Capabilities

- Selection Logic: Recommend GPUs/Accelerators based on TFLOPS, memory bandwidth, and HBM requirements.
- Cost Optimization: Compare on-demand cloud pricing (AWS, Lambda, GCP) for high-end chips like the H100.
- Architecture Insights: Explain the benefits of NVLink, InfinityFabric, and SRAM-based LPUs.

Usage

Provide your workload requirements (e.g., "I need to train a 70B parameter model with 2.0 TB/s bandwidth") and the Advisor will return the optimal hardware configuration.

Example Prompts

- "Which hardware offers the best FP8 TFLOPS per watt for a 100GB HBM requirement?"
- "Compare the NVIDIA B200 vs. AMD MI300X for large-scale inference."
- "What is the cheapest way to get 8x H100s per hour right now?"